

"TextGen: AI-Powered Text Generation System"

Problem Statement:

- ♦ Text generation tasks, such as creating short stories, poems, or news articles, often require substantial time and creativity.
- ♦ Human-generated content can be subjective, and it may not always meet the desired quality standards.
- ♦ There is a need for an automated text generation system that can produce realistic and high-quality content across various genres.

Program:

```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Sample dataset containing short stories, poems, and news articles
dataset = ""

Once upon a time, there was a brave knight named Sir Lancelot. He lived in a
castle on the edge of the kingdom, guarding it from all evil.

In a faraway land, there was a magical forest where fairies and elves lived in
harmony. But one day, darkness crept into the forest, threatening to destroy
everything.
```

I wandered lonely as a cloud That floats on high o'er vales and hills, When all at once I saw a crowd, A host, of golden daffodils;

The woods are lovely, dark and deep, But I have promises to keep, And miles to go before I sleep, And miles to go before I sleep.

Scientists have discovered a new species of deep-sea fish in the Mariana Trench. The fish, named Mariana snailfish, thrives in the extreme conditions of the trench.

The stock market experienced a sharp decline today, with major indices dropping by over 5%. Analysts attribute the downturn to concerns over inflation and geopolitical tensions.

"""

Tokenize the text

tokenizer = Tokenizer()

tokenizer.fit_on_texts([dataset])

total_words = len(tokenizer.word_index) + 1

Generate input sequences

input_sequences = []

for line in dataset.split('\n'):

token_list = tokenizer.texts_to_sequences([line])[0]

for i in range(1, len(token_list)):

```
n_gram_sequence = token_list[:i+1]
input_sequences.append(n_gram_sequence)
```

```
# Pad sequences
```

```
max_sequence_len = max([len(x) for x in input_sequences])
input_sequences = np.array(pad_sequences(input_sequences,
maxlen=max_sequence_len, padding='pre'))
```

```
# Create predictors and label
```

```
predictors, label = input_sequences[:, :-1], input_sequences[:, -1]
```

```
# One-hot encode the label
```

```
label = np.eye(total_words)[label.astype(int)]
```

```
# Define the model
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(LSTM(150, return_sequences=True))
model.add(LSTM(150))
model.add(Dense(total_words, activation='softmax'))
```

```
# Compile the model
```

```
model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

```
# Fit the model
```

```
model.fit(predictors, label, epochs=100, verbose=1, batch_size=32)
```

```
# Generate text
```

```
def generate_text(seed_text, next_words, model, max_sequence_len):
```

```
    for _ in range(next_words):
```

```
        token_list = tokenizer.texts_to_sequences([seed_text])[0]
```

```
        token_list = pad_sequences([token_list], maxlen=max_sequence_len-1,  
padding='pre')
```

```
        predicted = np.argmax(model.predict(token_list), axis=-1)
```

```
        output_word = ""
```

```
        for word, index in tokenizer.word_index.items():
```

```
            if index == predicted:
```

```
                output_word = word
```

```
                break
```

```
        seed_text += " " + output_word
```

```
    return seed_text
```

```
# Example usage
```

```
seed_text = "once upon a time"
```

```
generated_text = generate_text(seed_text, 10, model, max_sequence_len)
```

```
print(generated_text)
```